

✓ Hyperparameter Tuning for kNN for Predicting Heart Disease

1. Import "heart.csv".

2. Import Library

```
import pandas as pd
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import GridSearchCV
```

3. Load the Dataset into a frame

```
df = pd.read_csv('/content/heart (1).csv')
```

4. Print the description, dimensions and first five records of the frame.

```
print(df.info())
print(df.shape)
print(df.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
None
(303, 14)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

5. Check whether the data has any missing value in any column.

```
df.isnull().sum()
```

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```


6. Check whether the data has balanced class distribution. Class target = 0 indicates "Heart Disease" and target = 1 indicates "No Heart Disease".

```
df['target'].value_counts()
```

```
1    165
0    138
Name: target, dtype: int64
```

```
sns.countplot(x='target',data=df)
```

```
<Axes: xlabel='target', ylabel='count'>
```



7. Create input features X, target Y, classifier object, train-test-split using 80-20% split

```
x = df.drop(columns=['target'])
y = df['target']
```

```
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=4)
```

8. Train model

```
#Training the model
knn = KNeighborsClassifier()
knn.fit(x_train, y_train)
```

```
▸ KNeighborsClassifier
KNeighborsClassifier()
```

9. Validate model on test set

```
y_pred = knn.predict(x_test)
```

10. Print Classification Report on test data

```
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.48	0.52	0.50	25
1	0.65	0.61	0.63	36
accuracy			0.57	61
macro avg	0.56	0.57	0.56	61
weighted avg	0.58	0.57	0.58	61

11. Print AUC score on test data

```
roc_auc_score(y_test, y_pred)

0.5655555555555556
```

The performance of the model is very poor. Hence hyperparameters of kNN to be tuned using GridSearchCV.

12. Hyperparameter tuning using GridSearchCV. Set the parameters a) leaf-size= 1 to 15, b) n_neighbors = 1 to 10 and c) distance metric, p = 1, 2. When p = 1 its Manhattan and p = 2 its Euclidean distance. GridSearchCV uses CV to search for the optimal values of the hyperparameters. It accepts the hyperparameters as a dictionary.

```
leaf_size = list(range(1,15))
n_neighbors = list(range(1,10))
p=[1,2]
hyperparameters = dict(leaf_size=leaf_size, n_neighbors=n_neighbors, p=p)
```

13. Train a new kNN model using GridSearchCV.

```
knn_2 = KNeighborsClassifier()
clf = GridSearchCV(knn_2, hyperparameters, cv=10, scoring = 'roc_auc')
best_model = clf.fit(x,y)
```

14. Print the best values of the hyperparameters.

```
#Nilai hyperpaameters terbaik
print('Best leaf_size:', best_model.best_estimator_.get_params()['leaf_size'])
print('Best p:', best_model.best_estimator_.get_params()['p'])
print('Best n_neighbors:', best_model.best_estimator_.get_params()['n_neighbors'])
print('Best Score:', best_model.best_score_)

Best leaf_size: 9
Best p: 1
Best n_neighbors: 7
Best Score: 0.7483536683904332
```

15. Validate the model on test data

```
y_pred = best_model.predict(x_test)
```

16. Print classification report and AUC score of the model on test data

```
print(classification_report(y_test, y_pred))
print("AUC SCORE is",roc_auc_score(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.72	0.72	0.72	25
1	0.81	0.81	0.81	36
accuracy			0.77	61
macro avg	0.76	0.76	0.76	61
weighted avg	0.77	0.77	0.77	61

```
AUC SCORE is 0.7627777777777778
```